# The Role of GPUs and CPUs in Achieving Artificial General Intelligence via Symbolic AI

### 1. Introduction

Artificial General Intelligence (AGI) represents a significant theoretical milestone in the field of artificial intelligence, envisioning intelligent agents capable of understanding and executing any intellectual task that a human being can. This encompasses a broad spectrum of cognitive abilities, including abstract reasoning, emotional understanding, and the capacity for context-dependent learning, far exceeding the capabilities of current narrow AI systems designed for specific tasks. Among the various approaches proposed to achieve AGI, symbolic AI stands out as a paradigm that emphasizes the use of high-level symbolic representations of problems, employing logic, rules, and structured knowledge to perform complex tasks. Historically, symbolic AI was a dominant approach in the early stages of AI research, focusing on encoding human knowledge and behavioral rules into computer code. This method contrasts with the more recent surge in connectionist AI, which relies on learning patterns from vast amounts of data through neural networks. The fundamental difference in their approach suggests that the hardware best suited for each paradigm might also differ. This report aims to address the core question of whether AGI is likely to occur in the near future through symbolic AI, specifically focusing on the potential roles of Graphics Processing Units (GPUs) and Central Processing Units (CPUs) in driving this endeavor. The analysis will delve into the computational characteristics of symbolic AI, the architectural strengths and limitations of CPUs and GPUs, the broader landscape of AI hardware, and expert perspectives on the hardware requirements for AGI.

### 2. Understanding Symbolic AI and its Computational Needs

The foundation of symbolic AI lies in the explicit representation of knowledge using symbols, rules, and logical structures. This paradigm operates on the principle that intelligence can emerge from the manipulation of these symbolic representations based on predefined rules and logical inference. Various techniques are employed to achieve this, including logic programming languages like Prolog, which represent knowledge as logical statements and rules, semantic networks that depict concepts as nodes and their relationships as labeled links, frame-based systems that organize knowledge into data structures with slots for attributes and values, and production rules that express knowledge as condition-action pairs. The central process in symbolic AI involves manipulating these symbolic structures according to established rules to perform reasoning and problem-solving. This often entails computational operations such as logical inference, where new knowledge is derived from existing information through deductive, inductive, or abductive reasoning. Pattern matching and rule application are also crucial, where algorithms identify patterns in the symbolic representations and apply relevant rules to generate new information or take actions. For systems using semantic networks, graph traversal and manipulation are essential for navigating and extracting relationships between concepts. Furthermore, symbolic AI systems require mechanisms for querying and updating their knowledge bases to maintain and evolve their understanding of the domain. These computational operations often exhibit complex control flow, involving conditional branching and potentially irregular memory access patterns, characteristics that can influence the suitability of different hardware architectures.

Historically, the computational demands of symbolic AI led to the development of specialized hardware like LISP machines, which were designed to accelerate the development and execution of symbolic programs [1], S_21]. However, as knowledge bases in symbolic AI systems grow in size and complexity, scalability becomes a significant challenge. Moreover, symbolic AI has faced difficulties in handling the inherent uncertainty and ambiguity present in real-world information, as it typically relies on precise and unambiguous representations of knowledge. Compared to machine learning approaches, traditional symbolic AI systems often exhibit limited adaptability and learning capabilities, typically requiring manual programming and struggling to learn from new data without explicit reprogramming. The historical reliance on specialized hardware suggests that standard CPUs might not always be the optimal choice for computationally intensive symbolic AI tasks. Nevertheless, the later adoption of general-purpose Unix workstations for running LISP and Prolog indicates that advancements in processor technology and compiler optimization played a role in making symbolic AI more accessible on conventional hardware [1]. However, the fundamental challenges associated with scaling symbolic AI and effectively dealing with the complexities of real-world scenarios might constrain its potential as the primary driver of near-future AGI, regardless of the underlying hardware platform.

### 3. The Strengths and Limitations of CPUs for Symbolic AI

Central Processing Units (CPUs) are the traditional workhorses of general-purpose computing, and their architecture reflects this versatility. CPUs are designed with a focus on sequential processing, excelling at executing complex, single-threaded tasks with low latency. Their architecture incorporates strong control units that are highly adept at managing intricate logic and control flow, which is crucial for navigating the complex decision pathways often found in symbolic AI algorithms. Furthermore, CPUs feature large and sophisticated cache hierarchies, enabling fast access to frequently used data and instructions, a characteristic that can benefit the iterative and rule-based nature of symbolic processing. The very design of a CPU, optimized for a wide range of computational tasks, aligns reasonably well with the fundamental principles of symbolic AI, particularly the logical inference and rule-based processing that form its core. Many symbolic AI operations involve traversing complex decision trees and applying rules in a sequential manner, tasks that CPUs are inherently designed to handle efficiently. The historical use of general-purpose workstations powered by CPUs for running symbolic AI languages like LISP and Prolog further demonstrates the capability of CPUs to execute symbolic computations [1]. This indicates that CPUs have indeed been a viable platform for symbolic AI, although the performance might have been a limiting factor for very large and intricate systems.

However, CPUs also have limitations when considering the demands of potentially large-scale symbolic AI aimed at achieving AGI. Their primary constraint lies in their limited parallel processing capabilities compared to GPUs. While modern CPUs feature multiple cores, the sheer number of parallel processing units is significantly lower than that found in GPUs. This could pose a bottleneck if certain aspects of symbolic AI, such as the manipulation of very large knowledge bases or the potential for parallel rule evaluation, could benefit from a more parallel architecture. Although traditional symbolic AI has often been characterized by its sequential nature, the increasing scale and complexity required for AGI might necessitate exploring avenues for parallelization, where the inherent limitations of CPU architecture could become more pronounced.

### 4. The Potential of GPUs for Symbolic AI

Graphics Processing Units (GPUs) have emerged as a dominant force in modern AI, primarily due to their massively parallel architecture. Featuring thousands of cores designed for high throughput, GPUs excel at performing the same operation on multiple data points simultaneously, making them particularly well-suited for the parallel processing of large datasets, especially the matrix operations that are fundamental to deep learning. Modern GPUs even incorporate specialized cores like Tensor Cores, which are specifically designed to accelerate the matrix mathematics that underpin neural networks. This inherent strength in parallel computation presents an interesting contrast to the often sequential nature of traditional symbolic AI. Directly applying GPUs to purely symbolic AI tasks can be challenging because the irregular control flow and data dependencies common in many symbolic algorithms do not map efficiently onto the Single Instruction, Multiple Data (SIMD) architecture of GPUs. Furthermore, the overhead associated with transferring the symbolic data structures between the CPU and the GPU's memory can introduce performance bottlenecks.

Despite these challenges, there is growing interest in leveraging the immense computational power of GPUs for AI systems that incorporate symbolic reasoning. One promising direction is the field of neuro-symbolic AI, which seeks to integrate the strengths of neural networks (which run very efficiently on GPUs) with symbolic components to achieve more robust and interpretable AI systems. In such hybrid approaches, GPUs can be effectively used to accelerate the neural network components responsible for tasks like perception or learning, while symbolic processing for reasoning and decision-making might still occur on CPUs or potentially specialized hardware. Research in this area explores various ways to map symbolic AI operations onto parallel GPU architectures, such as parallel rule matching or graph processing for semantic networks. For instance, Large Language Models (LLMs), which are trained on vast datasets using GPUs, demonstrate an ability to handle human language as a form of symbols, suggesting a potential convergence between connectionist and symbolic paradigms. Therefore, while GPUs might not be a natural fit for all aspects of traditional symbolic AI, they could play a significant role in future AI systems that integrate symbolic reasoning, particularly within neuro-symbolic frameworks.

### 5. The Broader Landscape of AI Hardware for AGI

Beyond CPUs and GPUs, the field of AI hardware is rapidly evolving, with the emergence of various specialized accelerators designed to optimize specific AI workloads. Tensor Processing Units (TPUs), developed by Google, are specifically designed for the matrix operations prevalent in deep learning, offering high performance and energy efficiency for these tasks. Neural Processing Units (NPUs) are another type of specialized accelerator tailored for neural network computations, aiming to mimic the human brain's processing of information and offering efficiency gains in areas like image recognition and natural language processing. Field-Programmable Gate Arrays (FPGAs) provide a different approach, offering reconfigurable hardware that can be programmed to fit the specific needs of various AI tasks, enabling updates and modifications without hardware replacement. Application-Specific Integrated Circuits

(ASICs) are custom-designed chips optimized for very specific AI functions, offering superior performance and energy efficiency for those particular applications. The development of this diverse range of specialized hardware underscores a trend towards optimizing computational resources for different AI paradigms and tasks.

These specialized accelerators also hold relevance for neuro-symbolic approaches. For instance, TPUs and high-end GPUs can efficiently handle the computationally intensive neural network components of such systems. Furthermore, the reconfigurability of FPGAs could potentially be leveraged to implement custom circuits for accelerating specific symbolic logic operations. This suggests that future AI systems incorporating symbolic reasoning might benefit from heterogeneous hardware architectures, where different components are optimized for neural and symbolic processing. Looking further into the future, emerging technologies like quantum computing hold the potential to revolutionize AI hardware by offering unparalleled computational capabilities for certain complex AI problems. Additionally, there is a growing recognition of the need for more energy-efficient hardware to support the immense computational demands of advanced AI, including AGI. The long-term trajectory of AGI hardware might therefore involve entirely new computing paradigms or significant breakthroughs in energy efficiency that go beyond the current capabilities of CPUs and GPUs.

**6. Expert Perspectives on AGI Hardware**

Expert opinions on the hardware requirements for achieving AGI vary considerably, reflecting the inherent uncertainty surrounding the nature and computational demands of true general intelligence. Some perspectives suggest that the computational power needed for AGI using current machine learning methods could be astronomically high, potentially exceeding the planet's available resources. This viewpoint often leads to the suggestion that entirely new hardware paradigms, such as quantum computing, and significantly more energy-efficient solutions will be necessary. Conversely, other experts propose that human-level AGI might be achievable with surprisingly modest hardware, even suggesting that a current high-end gaming PC could potentially suffice. Considerations around the minimum hardware required for an AGI to become self-sufficient or pose a risk also contribute to this debate, with some analyses suggesting that very limited computational resources might not be sufficient for complex tasks like fluent language or sophisticated environmental simulation. Despite these differing views, there is a general consensus among leading AI companies that the path towards AGI will necessitate a substantial increase in AI infrastructure, leading to significant investments in AI hardware.

Regarding the specific hardware for symbolic AI, historical context reveals the use of specialized LISP machines, which were eventually superseded by general-purpose workstations as processor technology advanced [1]. More recent research in neuro-symbolic AI indicates that current off-the-shelf hardware, including CPUs and GPUs, often suffers from inefficiencies when running these hybrid models, particularly due to the memory-bound nature of symbolic operations and complex control flow. This has led to suggestions for the development of hardware acceleration specifically tailored for vector-symbolic architectures to improve the performance, efficiency, and scalability of neuro-symbolic computing. These expert observations collectively suggest that while general-purpose hardware can be used for symbolic AI, there is a recognized need for more optimized hardware solutions, especially for advanced neuro-symbolic approaches that aim to integrate symbolic reasoning with the computational power of neural networks.

## 7. Challenges and the Near-Term Outlook

Achieving AGI through symbolic AI faces significant challenges that extend beyond just the underlying hardware. Rule-based systems in symbolic AI can be brittle when confronted with novel situations not explicitly covered by their programmed knowledge. Encoding the vast amount of common-sense knowledge and handling the inherent complexities and nuances of the real world remain substantial hurdles for purely symbolic approaches. Furthermore, the limited ability of traditional symbolic AI to learn and adapt from data without explicit reprogramming contrasts sharply with the capabilities of modern machine learning. These inherent limitations of the symbolic AI paradigm itself might impede its progress towards AGI in the near future, irrespective of the computational platform employed.

From a hardware perspective, current off-the-shelf solutions might not be ideally suited for the computational characteristics of complex symbolic AI, particularly in neuro-symbolic settings. While CPUs excel at the sequential and logic-driven operations common in symbolic AI, they might struggle with potential avenues for parallelization that could arise in very large systems. Conversely, GPUs, while offering immense parallel processing power, might not be efficiently utilized by traditional symbolic AI due to architectural mismatches with the irregular control flow and data dependencies often found in symbolic algorithms. Specialized hardware designed for neuro-symbolic AI is still in the early stages of development, and its widespread availability and effectiveness remain to be seen.

Considering these limitations of both the symbolic AI paradigm and the current hardware landscape, it appears less likely that purely symbolic AI, driven solely by either CPUs or GPUs, will lead to AGI in the near future. Neuro-symbolic approaches, which aim to combine the strengths of both symbolic and connectionist methods and potentially leverage a combination of CPUs, GPUs, and specialized hardware, represent a more promising direction for incorporating symbolic reasoning into advanced AI systems. However, achieving true AGI through this hybrid route is also likely a long-term endeavor that extends beyond the "near future." The current trajectory of AI research and development heavily favors connectionist approaches, particularly deep learning, which have demonstrated remarkable progress in various domains. While these approaches also face significant challenges on the path to AGI, they currently appear to be the dominant focus of efforts aimed at achieving more general forms of artificial intelligence.

## 8. Conclusion

Based on the analysis, it is improbable that Artificial General Intelligence will emerge in the near future through the exclusive use of GPUs or CPUs for traditional symbolic AI. While CPUs possess architectural characteristics suitable for the sequential and logic-driven computations prevalent in symbolic AI, their inherent limitations in parallel processing might hinder scalability for highly complex AGI systems. GPUs, despite their massive parallel processing capabilities that have revolutionized connectionist AI, do not align as naturally with the typically sequential and logic-intensive operations of traditional symbolic AI, potentially leading to inefficient utilization of their resources.

Looking ahead, neuro-symbolic approaches, which seek to integrate the strengths of both neural and symbolic methods, hold greater promise for incorporating symbolic reasoning into advanced AI systems. These hybrid models might leverage a combination of CPUs, GPUs, and potentially specialized hardware to accelerate different aspects of the computation. However, achieving true AGI through this route remains a significant long-term challenge. The development of novel hardware architectures specifically designed to address the unique computational demands of symbolic or neuro-symbolic computation could be crucial for future progress in this area. In the immediate future, advancements in connectionist AI continue to be the primary focus in the pursuit of more general artificial intelligence, although the path to true AGI remains complex and uncertain, regardless of the specific AI paradigm or underlying hardware platform.

**Table 1: Comparison of CPU and GPU Architectures for AI Tasks**

| Feature | CPU | GPU |
|---|---|---|
| Processing Paradigm | Sequential | Parallel |
| Core Count | Few (tens to hundreds) | Many (thousands) |
| Core Complexity | Complex, designed for a wide range of tasks | Simpler, optimized for parallel throughput |
| Memory Architecture | Sophisticated multi-level cache hierarchy for low latency | Focus on high memory bandwidth for parallel data access |
| Primary Strengths | Low latency for single-threaded tasks, strong control units, versatile for general-purpose computing | High throughput for parallel tasks, excellent for data-parallel computations |
| Primary Weaknesses | Limited parallel processing capability | Higher latency for sequential tasks, less efficient for complex control flow |
| Typical AI Use Cases | Control tasks, data preprocessing, some smaller AI models | Training and inference of deep learning models, parallelizable AI algorithms |

| | | |
|---|---|---|
| Suitability for Traditional Symbolic AI | Moderate; well-suited for sequential logic and rule-based systems, but potential scalability limitations | Low; architectural mismatch with sequential and irregular control flow |
| Suitability for Neuro-Symbolic AI | Can handle control and potentially some symbolic components | Excellent for accelerating neural network components |

**Table 2: Specialized AI Hardware and Their Applications**

| **Hardware Type** | **Primary Architecture Characteristics** | **Key Strengths** | **Typical AI Use Cases** | **Relevance to Neuro-Symbolic AI** |
|---|---|---|---|---|
| TPU (Tensor Processing Unit) | Custom ASIC designed for matrix operations | High performance and energy efficiency for tensor computations | Deep learning training and inference, especially with TensorFlow | Excellent for accelerating neural network components |
| NPU (Neural Processing Unit) | Architecture optimized for neural network operations | Efficiency for specific AI tasks like image recognition and NLP | Edge AI applications, mobile devices | Potential for accelerating specific neural components |
| FPGA (Field-Programmable Gate Array) | Reconfigurable integrated circuits | Flexibility, low latency, customizable hardware | Real-time processing, AI inference at the edge, prototyping new AI algorithms | Potential for implementing custom symbolic logic circuits |
| ASIC | Custom-design | Highest | Specific AI | Potential for |

| | | | | |
|---|---|---|---|---|
| (Application-Specific Integrated Circuit) | ed chips for specific tasks | performance and energy efficiency for targeted applications | inference tasks, high-volume deployments | highly optimized symbolic or neural components |

**Table 3: Advantages and Limitations of Symbolic AI**

| **Aspect** | **Advantages** | **Limitations** |
|---|---|---|
| Knowledge Representation | Explicit and human-readable, precise modeling of complex relationships | Time-consuming and laborious knowledge encoding |
| Reasoning | Formal logic allows for clear reasoning and deduction, interpretable decision-making | Struggles with common-sense reasoning |
| Data Requirements | Does not require vast amounts of training data | Requires well-defined and structured knowledge |
| Adaptability | Suitable for well-defined domains | Challenging to handle unstructured data, lacks self-learning capabilities |
| Handling Uncertainty | Can incorporate logical rules for handling some uncertainty | Difficulty with ambiguous or incomplete information |
| Scalability | Effective for smaller, well-defined domains | Scalability issues with large knowledge bases and complex rule sets |

## Works cited


1. Symbolic artificial intelligence - Wikipedia, accessed March 14, 2025, https://en.wikipedia.org/wiki/Symbolic_artificial_intelligence

2. Symbolic AI: Revolutionizing Rule-Based Systems - SmythOS, accessed March 14, 2025, https://smythos.com/ai-agents/ai-tutorials/symbolic-ai-applications/
3. How Close Are We to AGI and What Stands in the Way? · Neil Sahota, accessed March 14, 2025, https://www.neilsahota.com/how-close-are-we-to-agi-and-what-stands-in-the-way/